

Supplementary Material of CloudDet

I. QUANTITATIVE EVALUATION

To evaluate the performance and the scalability of the proposed anomaly detection algorithm, we tested our proposed algorithm through a quantitative comparison with baseline methods based on Yahoo! S5 real time-series dataset [9].

A. Dataset, Baselines and Evaluation Metrics.

Dataset. In our evaluation, we used the A1Benchmark file of Yahoo! S5-A Labeled Anomaly Detection dataset, which contains 67 real time series with labeled anomalies. These real time series were collected from the real production traffic to some of the Yahoo! properties. Note that the timestamps of the A1Benchmark were replaced by integers with the increment of 1, where each data-point represents 1 hour worth of data.

Baseline Methods and its Parameter Settings. In order to decide which anomaly detection algorithms should be compared, two principles were applied: (1) cover typical anomaly detection techniques from different categories; and (2) the execution time is not much longer than our proposed algorithm. By surveying the existing paper, we chose five representative anomaly detection algorithms from five categories:

One-Class Support Vector Machine (oc-SVM), from classification-based algorithms, uses a hyperplane to distinguish two classes [4]. *RBF* (radial basis function) kernel is used in our system to deal with high-dimensional data. The kernel coefficient γ is chosen as the adjustable parameter for this algorithm when conducting the accuracy evaluation, which can grow proportionally from 0.001 to 0.01, with 0.001 as the step length.

Local Outlier Factor (LOF) is also one of the neighbor-based analysis methods, but it is density-based [2]. It determines an outlier instance a by comparing a 's k -neighborhood density to the k -neighborhood density of a 's k -neighbors. We select k as the adjustable parameter for LOF's accuracy evaluation, which contains ten values growing proportionally from 20 to 38, with 2 as the step length.

Robust Covariance Estimation (RCov) is a statistic-based algorithm that assumes the data follow a known distribution (e.g., Gaussian distribution). We use the Mahabolis distances to determine the outlyingness of a point from the known distribution. We set the proportion of points to be included in the support of the raw MCD (minimum covariance determinant) estimate, s , as the parameter for this algorithm's accuracy evaluation. It can grow proportionally from 0.05 to 0.95, with 0.1 as the step length.

Bitmap Detector (BD), based on symbolic aggregate approximation (SAX) of time series, is an assumption-free anomaly detection algorithm in time series via the bitmap [8]. We select the number of sections to categorize values as the adjustable

parameter, which grows from proportionally from 2 to 20, with 2 as the step length.

Twitter Anomaly Detection (Tad), also referred to as Seasonal Hybrid ESD (S-H-ESD), builds upon the Generalized ESD test for detecting anomalies. The algorithm employs piecewise approximation – this is rooted to the fact that trend extraction in the presence of anomalies in non-trivial – for anomaly detection. The periodic basis (the proportion of the time series length) is selected as the adjustable parameter, which changes from proportionally from 0.03 to 0.3, with 0.03 as the step length.

We also tried Isolation Forest (iForest) [7] from model-based method, and the HOT-SAX [6] from SAX-based method for comparisons. However, we decided not to show their results due to their extremely long execution times compared with our proposed algorithms.

Evaluation Metrics. The evaluation metrics (ROC and execution time) are the standard information retrieval metrics used in many works related with time series anomaly detection [1], [5], [3]. We chose ROC for the accuracy evaluation as the ratio between positive and negative instances was extremely imbalanced. We chose execution time to evaluate the scalability of our algorithm.

B. Implementations and Evaluation Settings

Implementations. We used the implementations of Rcov, LOF and SVM in the scikit-learn package¹, the implementation of BD in the luminol package², and the Tad implementation in the pyculiarity (A Python port of Twitter's AnomalyDetection R Package)³. All the parameters were set by default, except for the adjustable parameter for each algorithm mentioned above.

Accuracy. First, we tested the accuracy of each algorithm with ten different values of its chosen parameter (refer to Baseline Methods above). In particular, our proposed algorithm took the number of recent "historical data" L as the adjustable parameter (see Section 4.1 in paper), which grew from 5 to 50, with 5 as the step length. Then we ran each parameter setting of the algorithm ten times, with each time we choosing the discrimination threshold (proportion of anomalies) [0.005, 0.01, 0.02, 0.04, 0.08, 0.16, 0.32, 0.64, 0.8, 0.95]. The thresholds grew exponentially, except for the last two. Therefore, we got 100 runs (10 parameter values \times 10 discrimination thresholds) for each algorithm in the accuracy evaluation, which reduced the bias of these algorithms to different datasets.

¹https://scikit-learn.org/stable/modules/outlier_detection.html

²<https://github.com/linkedin/luminol>

³<https://github.com/zrnsmpyculiarity>

To verify that our proposed algorithm is able to produce satisfactory detection of time series anomalies with different patterns, we benchmark the algorithm on a set of datasets selected from AIBenchmark. The specific datasets we used include: real_5, real_31 and real_32. The detailed results are described as follows.

(1) Fig. 1(1) displays the accuracy testing for six anomaly detection algorithms with the real_5 dataset based on ROC curve. Our algorithm is denoted as “AS”. The dataset contains 2 anomalies in 1439 points, which can be recognized as short-term spike anomalies in the time series. The ROC plot shows that most algorithms have a good performance for detecting the short spikes in the data except Tad. Specifically, our algorithm, RCov and BD have a higher true positive rates than the others when the false positive rates remain low (below 0.1). The reason that all of them have a sound performance might be that this kind of anomaly pattern is common and simple in data distribution for anomaly detection.

(2) Fig. 1(2) shows the ROC curves for different algorithms with the real_31, which contains mixed change patterns (both long-term and short-term spike anomalies) with a 1.68% outlier percentage (24/1427). Overall, our algorithm (denoted as “AS”) outperforms the five baseline methods when there are different anomaly patterns in data. In particular, our algorithm had a higher true positive rates when the false positive rate was low (below 0.2), which means that our algorithm can reduce false positive anomalies when detecting the same number of anomalies as the baseline algorithms. Different from Fig. 1(1), the LOF and Tad also have a sound performance while the BD and SVM have the worst performance. This result indicates the superiority of our algorithm when there are diverse patterns in time-series data. This is important for cloud computing anomaly detection because the experts can save efforts in anomaly diagnosis when the data scale is very large and variant.

(3) Fig. 1(3) shows the accuracy testing results for six anomaly detection algorithms with the real_32 dataset. There are 37 outlier points in 1427 observations, which are two long-term anomalous periods. The ROC plots show that our algorithm (“AS”), have a comparable performance with LOF, RCov and Tad. By contrast, BD suffers from finding more anomalies with the same false positive rates, which means it might not be specialized for long-term anomalies.

In general, our proposed algorithm are better than other baseline approaches based on the ROC plots, especially when the time series contain different types of anomaly patterns.

Scalability. The proposed algorithm must be scale-out and efficient to cope with large-scale data generated by cloud computing systems. To perform the evaluation, we tested the execution time of each algorithm by varying the length of the input time-series data. Specifically, in the experiments, the algorithm was executed on 50 time-series datasets from the AIBenchmark (real_1 to real_50), with each dataset running ten times with its ten adjustable parameter values. We only varied the length of the input time series in different experiments to determine the execution time of the algorithm. We varied

the length of time series, from 100 to 700 points with 100 as the step length, by selecting the first 100–700 data points of a dataset. The experiments were conducted in an eight-core (Intel Core i7-4790 CPU@3.6GHz) Window 10 computer with 16 GB memory. The results are summarized in Fig. 1(4). The figure suggests that the execution time of our algorithm can scale and exhibits linearly with the length of the time series. “AS” is less scalable than RCov and LOF, probably due to the time spent for pattern matching before anomaly detection, but the difference is acceptable considering the higher accuracy in detection results compared with other algorithms. Although Tad performs well in accuracy test in general, it has a worst computing speed compared with others.

C. Limitations.

Although the results showed that our proposed algorithm had sound performance when considering the speed and accuracy together, the validity of the results needs further evaluation. There exist some factors that may affect the validity: (1) The data may contains bias in anomaly patterns that are specialized for some anomaly detection algorithms. For example, LOF has a good performance for data with long-term anomalies ,as shown in Fig. 1(3). Therefore, more datasets could be used to conduct the evaluation as future work. Moreover, the results of the proposed algorithm could be affected by the aggregation strategy of the anomaly scores from the three components (i.s., spike, period and trend). A more optimal or automated aggregation method should be proposed in the future work; (2) The results could also be affected by the parameter selections. There is only one adjustable parameter in the current evaluation, which might induce the bias of the evaluation results; (3) We need more evaluation metrics to conduct a comprehensive evaluation of our proposed algorithm. The accuracy test could be extended with other metrics like precision-recall rate, and the scalability test can vary in other variables like the node numbers.

REFERENCES

- [1] B. Agrawal, T. Wiktorski, and C. Rong. Adaptive real-time anomaly detection in cloud infrastructures. *Concurrency and Computation: Practice and Experience*, 29(24):e4193, 2017.
- [2] M. M. Breunig, H.-P. Kriegel, R. T. Ng, and J. Sander. Lof: Identifying density-based local outliers. In *ACM SIGMOD Record*, vol. 29, pp. 93–104. ACM, 2000.
- [3] N. Cao, C. Lin, Q. Zhu, Y.-R. Lin, X. Teng, and X. Wen. Voila: Visual anomaly detection and monitoring with streaming spatiotemporal data. *IEEE Transactions on Visualization and Computer Graphics*, 24(1):23–33, 2018.
- [4] C.-C. Chang and C.-J. Lin. Libsvm: a library for support vector machines. *ACM Transactions on Intelligent Systems and Technology (TIST)*, 2(3):27, 2011.
- [5] S. Huang, C. Fung, C. Liu, S. Zhang, G. Wei, Z. Luan, and D. Qian. Arena: Adaptive real-time update anomaly prediction in cloud systems. In *2017 13th International Conference on Network and Service Management (CNSM)*, pp. 1–9. IEEE, 2017.
- [6] E. Keogh, J. Lin, and A. Fu. Hot sax: Finding the most unusual time series subsequence: Algorithms and applications. In *Proc. of the 5th IEEE Int’l. Conf. on Data Mining*, pp. 440–449. Citeseer, 2004.
- [7] F. T. Liu, K. M. Ting, and Z.-H. Zhou. Isolation forest. In *Eighth IEEE International Conference on Data Mining*, pp. 413–422. IEEE, 2008.

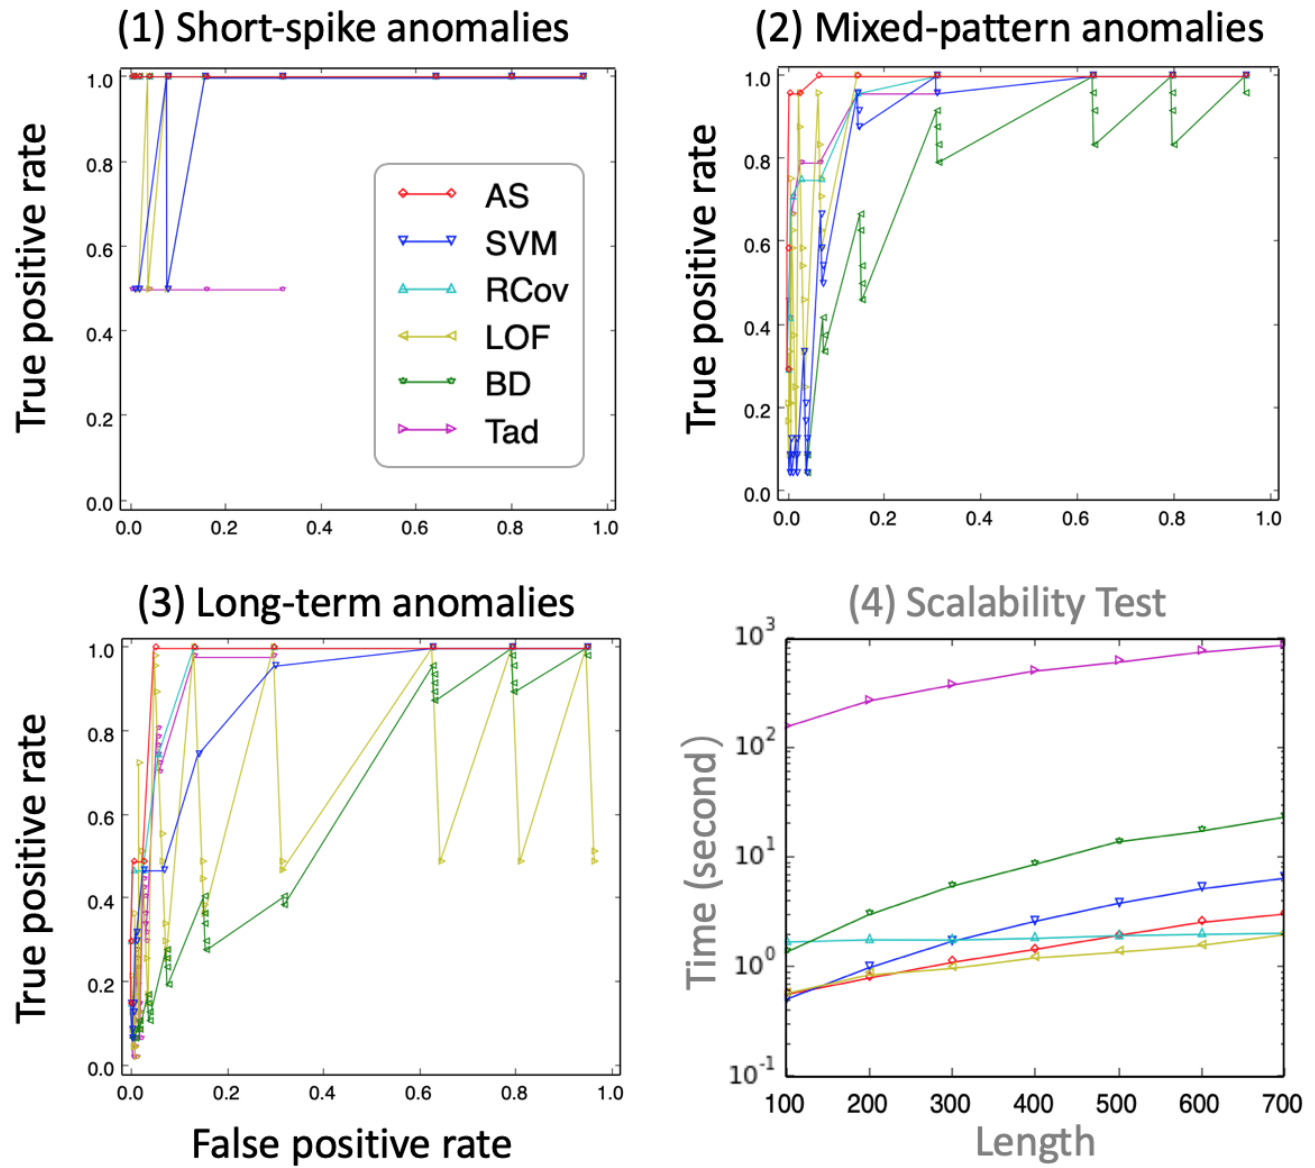


Fig. 1. Quantitative validation results. (1) – (3) shows the ROC curves of six algorithms based on three real-world datasets that contain the short-term spikes, mixed anomaly patterns and long-term spikes, respectively. The results indicate that our algorithm (“AS”) has a sound performance in accuracy with different anomaly patterns, especially in the mixed pattern. (4) shows the scalability of six algorithms. Our algorithm can scale and exhibit linearly with the varying length of time series.

[8] L. Wei, N. Kumar, V. N. Lolla, E. J. Keogh, S. Lonardi, and C. A. Ratanamahatana. Assumption-free anomaly detection in time series. In *SSDBM*, vol. 5, pp. 237–242, 2005.

[9] Yahoo. S5-a labeled anomaly detection dataset. ”<https://webscope.sandbox.yahoo.com/catalog.php?datatype=s&did=70>”.